

Causal Inference Crash Course:

Part 10 — Conformal Inference

Julian Hsu

Causal Inference Series

1. Foundations
2. Causal Models
3. Inference
4. Best Practices
5. Heterogeneous Treatment Effects
6. Panel Models
7. Regression Discontinuity
8. Surrogate Models
9. Arguable Validation
10. **Conformal Inference** ← **this deck**

Overview

- **So what?** Every model's point prediction is wrong by some amount — conformal inference is how you say honestly by how much. It turns *any* prediction into a **set** with a **finite-sample, distribution-free coverage guarantee**, needing only **exchangeability**.
- The payoff: this same machinery gives **valid causal inference with a single treated unit** (synthetic control), where classical standard errors go quiet.
- We cover:
 - the idea: **split conformal** and its coverage guarantee
 - the one assumption — **exchangeability** — and what breaks it
 - two fixes when it fails: **moving blocks** and **weighting**
 - the causal payoff: conformal inference for **synthetic controls**
- Assumes the earlier **Inference** and **Panel Models** decks.

The Idea

You don't have to trust the model — you just have to believe the order of the data doesn't matter.

What we want

- **So what?** Your model gives a point prediction. How often is it *this* far off? Most workflows just don't say — conformal inference answers that honestly, without pretending the model is right.
- Classical prediction intervals need the **model to be right** — Wald wants a correct parametric form, Bayes a correct likelihood. Neither survives contact with the black-box model you actually shipped.
- **Conformal** asks for far less: a set $\mathcal{C}(x)$ with

$$\mathbb{P}(Y_{n+1} \in \mathcal{C}(X_{n+1})) \geq 1 - \alpha$$

in **finite samples**, for **any** distribution, wrapped around **any** predictor — however ugly.

- Running example we'll carry through: predict a customer's next-quarter purchases, and give an interval that really covers 90% of the time.

The one assumption: exchangeability

- **So what?** Here's the trade: give up on knowing the model is right, and in exchange you only need to believe the data doesn't care about order. That's a much easier thing to defend.
- Formally: $(X_1, Y_1), \dots, (X_{n+1}, Y_{n+1})$ are **exchangeable** — their joint distribution is unchanged by reordering them. (i.i.d. is the familiar special case.)
- That's the *whole* assumption — no distributional form, no correctly specified model. Everything downstream is a consequence of it.
- And it's exactly what **fails** for time series and panels — the moment your data has a "before" and "after," this assumption is in trouble. We'll come back to that.

Conformity scores

- **So what?** Before you can say "how often are we wrong," you need a ruler for "how wrong." That ruler is the **nonconformity score**.
- Pick a score $s(x, y)$ — "how strange is this pair, to the model?" Larger means a worse fit.
- Regression default: the absolute residual of any fitted $\hat{f}$,


$$s(x, y) = |y - \hat{f}(x)|.$$

- Good news: the guarantee holds for **any** score and **any** model. A *good* score doesn't buy you validity — you already have that. It only makes the resulting set **tighter**.

Split conformal — the recipe

The mechanics, step by step. For miscoverage level α , split the data into a **training** fold and a **calibration** fold of size n :

1. **Fit** $\hat{f}$ on the training fold only.
2. **Score** the calibration set: $s_i = s(X_i, Y_i)$, $i = 1, \dots, n$.
3. **Quantile**: $\hat{q} = s_{(k)}$, the k -th smallest score, with $k = \lceil (n + 1)(1 - \alpha) \rceil$.
4. **Predict**: $\mathcal{C}(x) = \{y : s(x, y) \leq \hat{q}\}$.



The coverage guarantee


- **So what?** This is the payoff of the recipe on the last slide: a real, provable promise about how often your set is right — not "should be fine in practice."
- **Theorem** (Vovk et al. 2005; Lei et al. 2018). If the data are **exchangeable**, split conformal satisfies

$$1 - \alpha \leq \mathbb{P}(Y_{n+1} \in \mathcal{C}(X_{n+1})) \leq 1 - \alpha + \frac{1}{n+1}.$$

- **Finite-sample, distribution-free, model-free.** The upper bound needs only that scores have no ties.
- Best part: the model can be **arbitrarily wrong** — coverage still holds. A bad model just gets you an honestly **wide** set, not a broken guarantee.

Why it works

- Insert the unseen test score as an $(n+1)$ -th point. By exchangeability its **rank is uniform** on $\{1, \dots, n+1\}$, so it exceeds the k -th calibration value only about α of the time — exactly the miscoverage event.
- So coverage is a **property of the procedure**, not luck: over repeated draws it follows a $\text{Beta}(k, n+1-k)$ law centered at $1 - \alpha$.



The p-value view


- **So what?** This reframing is the bridge to everything that follows — once conformal is "which values would I fail to reject," you can bolt on hypothesis-testing machinery instead of reinventing quantiles.
- Equivalent framing (Shafer & Vovk 2008): conformal is **hypothesis testing turned into prediction**. For each candidate label y ,

$$p(y) = \frac{1 + \#\{i : s_i \geq s(x, y)\}}{n + 1}, \quad \mathcal{C}(x) = \{y : p(y) > \alpha\}.$$

- Keep the labels you **cannot reject** at level α .
- This is the view we'll lean on for synthetic control — where the "label" becomes a hypothesized treatment effect, and rejecting it is the whole game.

Adaptivity — and its limit

- **Heads up:** the guarantee we just proved is **marginal**, not conditional — some regions can be **under-covered** and others **over-covered**; exact conditional coverage is impossible distribution-free.
- Fix with a smarter score, not a new guarantee: normalize by a local scale $s = |y - \hat{f}(x)| / \hat{\sigma}(x)$, or use **conformalized quantile regression** (Romano, Patterson & Candès 2019).




When Exchangeability Fails

Real data has a "before" and "after." That breaks the one assumption we just leaned on — here's what to do about it.

Time and treatment break exchangeability

- **So what?** Shuffle the data and normally nothing changes. Time series don't play along — and most data we actually work with is a time series.
- Panels and time series are **serially dependent** and often **drift** — reorder them and the joint law changes, so exchangeability fails.
- Worse for causal work: the null holds only **before** treatment, and the post-period is exactly what we suspect differs.
- Plain conformal then **under-covers** — calibrated on a stale past that no longer resembles the present.



Fix A — permute blocks, not points

- **So what?** If shuffling the data breaks its story, don't shuffle points — shuffle in chunks big enough to keep the story intact.
- Keep **contiguous blocks** of residuals together and permute the block **positions** — the moving-block idea borrowed from the block bootstrap.
- Local dependence rides along *inside* each block; only its placement is randomized, so the null datasets keep the real time structure.
- **Exact** under exchangeability, **approximately valid** under stationarity and weak dependence — the block length must exceed the dependence range (Chernozhukov, Wüthrich & Zhu 2018, 2021).

Fix B — weight what belongs

- **So what?** The alternative move: don't chunk the data, just admit some points matter more than others — and say so explicitly, with weights fixed in advance.
- **Nonexchangeable conformal** (Barber, Candès, Ramdas & Tibshirani 2023): attach fixed weights w_i and take **weighted quantiles** of the scores — for time series, weight **recent** points more.
- The guarantee degrades gracefully rather than breaking outright:

$$\mathbb{P}(Y_{n+1} \in \mathcal{C}(X_{n+1})) \geq 1 - \alpha - \Delta,$$

where Δ is a weighted **distance from exchangeability**; it vanishes when the up-weighted points really do resemble the target.

- One rule: the weights must be **fixed a priori** — not learned from the scores, or you're back to cheating.

Two relaxations of one idea

- **Back to the big idea:** both fixes are the same move — restore just enough symmetry for the uniform-rank argument to go through, without pretending the data is something it isn't.
- **Blocks** — the *discrete* relaxation: glue dependent points together so a permutation can't shatter the local structure.
- **Weights** — the *continuous* relaxation: down-weight the points that don't belong with the target.

Conformal Inference for Synthetic Controls

The payoff: turn this same machinery into a causal test that works with just one treated unit.

The CWZ recast

- **So what — and here's the payoff we've been building to:** everything up to now was about intervals for predictions. This section turns the same machinery into a **causal test** — with just one treated unit, no fancy standard errors required.
- **Chernozhukov, Wüthrich & Zhu (2021)** glue two familiar problems together: causal inference = **counterfactual prediction** + **a test for a structural break**.
- Predict the treated unit's *untreated* path (synthetic control, DiD, factor models...), then ask whether the **post-treatment residuals break** from the pre-treatment ones.
- The pitch: **inference that doesn't care how good your synthetic control is** — validity comes from the permutation, not from a correct model.



Setup — we predict the counterfactual

- **So what?** Before the test, we need the ingredients: a treated unit, a donor pool, and a prediction of what would've happened without treatment — nothing new if you sat through the Panel Models deck.
- One treated unit Y_{1t} , a donor pool Y_{0t} , pre-periods $t \leq T_0$ and post-periods $t > T_0$. Fit a **synthetic control** $\hat{Y}_{1t}(0)$ (from the Panel Models deck).
- The **gap** $\hat{u}_t = Y_{1t} - \hat{Y}_{1t}(0)$ is the residual. Under *no effect*, that residual process has **no break** at T_0 .
- The residuals over time are serially dependent — so this is exactly the non-exchangeable case, handled by **moving-block** permutations from the last section.

Step 1 — impose the null, read the residuals

- **So what's actually being tested?** A hypothesized effect size θ — test a sharp null on the post-period effect, $H_0 : \tau_t = \theta$. Subtract it from the **post outcomes only**, refit over the **whole timeline**, and take residuals for every t :

$$\tilde{Y}_{1t} = Y_{1t} - \theta \cdot \mathbf{1}\{t > T_0\}, \quad \hat{u}_t = \tilde{Y}_{1t} - \hat{Y}_{1t}(0)$$



Step 2 — a statistic and a moving-block permutation


- **So what?** We need one number for "how weird do the post-period residuals look," and a null distribution to compare it against — that's what turns "residuals look big" into an actual p-value.
- Summarize the post-period residuals with a single number ($q = 1$ = mean absolute residual):

$$S_q(\hat{u}) = \left(\frac{1}{T_{post}} \sum_{t > T_0} |\hat{u}_t|^q \right)^{1/q}$$

- Compare S_q to its distribution under **moving-block permutations** of the residual sequence — the serial-dependence-robust scheme from the last section.
- `panelib`'s `conformal_inf` implements exactly this, so you don't have to.

Step 3 — the permutation p-value

$$p = \frac{1}{|\Pi|} \sum_{\pi \in \Pi} \mathbf{1}\{S_q(\hat{u}_\pi) \geq S_q(\hat{u})\}, \quad \text{reject } H_0 \text{ if } p \leq \alpha.$$




- **Careful:** this permutes residuals **across time** for one treated unit — not the Abadie placebo test, which permutes treatment **across units**.

Step 4 — confidence interval by test inversion

- **So what does this buy you?** A full confidence interval, not just a yes/no test — repeat the test over a grid of candidate effects θ and keep every value you **cannot reject**:

$$CI_{1-\alpha} = \{ \theta : p(\theta) > \alpha \}$$



- The point estimate sits where $p(\theta)$ peaks; the CI is the set of effects consistent with "no break."

Why this matters

- **The payoff, one more time:** this works with **one treated unit and few pre-periods** — and is **exact** under exchangeability, **robust** to serial dependence and to a **misspecified** counterfactual (validity comes from residual symmetry, not a correct model).
- The fine print: under dependence the guarantee is approximate, and it needs the counterfactual's **error to be stable over time** — a model whose bias *trends* into the post-period breaks a symmetry no permutation can restore.
- Resolution matters too: with very few post-periods the permutation p-value is coarse — in the Panel deck's simulations (2 post-periods) coverage ran **80–85%**, not 95%.
- Contrast Abadie placebo inference: its p-value resolution is only $1/(J + 1)$ and it needs a large donor pool.
- This is the inference the **Panel Models** deck used to put DiD, SC, and SDID on the same footing.

Conclusion

- **Back to the big idea:** conformal inference buys an **honest, finite-sample guarantee** from **one assumption you can actually reason about** — exchangeability — and it degrades gracefully, rather than catastrophically, when that assumption is stressed.
- Reach for it to **wrap black-box predictors** you don't fully trust, and — the reason we built this whole deck — to do inference where classical SEs go silent: **a single treated unit over time**.
- The price: coverage is **marginal**, not conditional; and under dependence you must **block** or **weight** to keep the guarantee.
- Bottom line: start from the guarantee you can actually defend, then pick the score and the permutation that fit your data.

Appendix

Appendix: split vs. full conformal

- **Split (inductive)** conformal — the recipe here — spends part of the data on a calibration fold. Cheap (one fit), but the interval is random in that split and a bit wider.
- **Full** conformal refits for every candidate label — exact but expensive; **CV+ / jackknife+** (Barber et al. 2021) recover most of the efficiency without a held-out split.
- Same guarantee; the choice is **data efficiency vs. compute**.

Appendix: the coverage law

- Conditional on the calibration draw, the coverage of split conformal is a **random variable**:

$$\mathbb{P}(Y_{n+1} \in \mathcal{C} \mid \text{calibration}) \sim \text{Beta}(k, n + 1 - k), \quad k = \lceil (n + 1)(1 - \alpha) \rceil.$$

- Its mean is $\geq 1 - \alpha$; the spread shrinks like $1/\sqrt{n}$ — the middle histogram made this concrete.
- The upper bound $1 - \alpha + \frac{1}{n+1}$ needs the scores to have **no ties** (continuous scores, or a random tie-break).

Appendix: another route — prediction intervals

- Conformal is not the only distribution-aware option for synthetic control.
- **Cattaneo, Feng & Titiunik (2021)** build **prediction intervals** that separate two error sources: the in-sample **weight-construction** error and the out-of-sample **stochastic** error, with finite-sample guarantees.
- Rule of thumb: **conformal** tests a sharp null by permutation and inverts to a CI; **prediction intervals** model the two error sources directly. Different assumptions, complementary diagnostics.

Appendix: conformal for treatment effects, beyond panels

- **Individual treatment effects** — Lei & Candès (2021): conformal intervals for $Y(1) - Y(0)$ under unconfoundedness, using *weighted* conformal to bridge the treated/control covariate shift. The natural companion to the HTE deck.
- **Covariate shift** — Tibshirani, Barber, Candès, Ramdas (2019): weighted conformal stays valid when the test point comes from a different covariate distribution with a known (or estimated) likelihood ratio.
- **Drift in production** — Gibbs & Candès (2021): *adaptive* conformal updates the level α_t online to keep realized coverage on target under distribution shift — the monitoring use case for deployed models.

Literature Review of Related Papers

- **Conformal prediction — foundations**

- Vovk, Gammerman, Shafer (2005). *Algorithmic Learning in a Random World*. Springer — the founding monograph. [link](#)
- Shafer & Vovk (2008). "A Tutorial on Conformal Prediction." *JMLR* 9, 371–421. [link](#)
- Lei, G'Sell, Rinaldo, Tibshirani, Wasserman (2018). "Distribution-Free Predictive Inference for Regression." *JASA* 113(523). [link](#)
- Angelopoulos & Bates (2023). "A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification." *Found. & Trends ML*. [link](#)

Literature Review (continued)

- **Adaptivity & beyond exchangeability**

- Romano, Patterson, Candès (2019). "Conformalized Quantile Regression." *NeurIPS*. [link](#)
- Barber, Candès, Ramdas, Tibshirani (2021). "Predictive Inference with the Jackknife+." *Ann. Statist.* 49(1). [link](#)
- Barber, Candès, Ramdas, Tibshirani (2023). "Conformal Prediction Beyond Exchangeability." *Ann. Statist.* 51(2), 816–845. [link](#)
- Tibshirani, Barber, Candès, Ramdas (2019). "Conformal Prediction Under Covariate Shift." *NeurIPS*. [link](#)
- Gibbs & Candès (2021). "Adaptive Conformal Inference Under Distribution Shift." *NeurIPS*. [link](#)

Literature Review (continued)

- **Conformal inference for causal panels / synthetic control**

- Chernozhukov, Wüthrich, Zhu (2018). "Conformal Inference Methods for Predictive ML with Dependent Data." *COLT (PMLR 75)*. [link](#)
- Chernozhukov, Wüthrich, Zhu (2021). "Conformal Inference for Counterfactual and Synthetic Controls." *JASA* 116(536). [link](#)
- Cattaneo, Feng, Titiunik (2021). "Prediction Intervals for Synthetic Control Methods." *JASA* 116(536). [link](#)
- Firpo & Possebom (2018). "Synthetic Control Method: Inference, Sensitivity Analysis and Confidence Sets." *J. Causal Inference* 6(2). [link](#)
- Lei & Candès (2021). "Conformal Inference of Counterfactuals and Individual Treatment Effects." *JRSS-B* 83(5). [link](#)

- **Companion code:** `panelib.py` (`conformal_inf`) in [statanomics](#).

